

Introduction to Web Services



Outline

- A bit of history:
 - RPC,
 - CORBA, and
 - RMI
- Web Services:
 - SOAP (very brief)

My slides have been supplemented by material from various sources, including: W3C, Lillian Cassel (Villanova), Chiyong Seo (USC), Alfred C. Weaver (UVa)



Communicating systems

- Researchers in the computing field have long worked on ways to **send and receive data** to between computers.
- **ARPANET**, a network connecting computers provided by the Advanced Research Project Agency (ARPA) has been introduced as one of the initial answers to this problem, in the 1960's.
- Its successor, the **Internet**, followed soon.




Communicating systems

- Then the question was: how to create software systems that could **execute programs remotely and receive their results?**
- And a related question: how to make implementing this remote calls between computers **easier**.
- As an answer, people have first introduced
 - **RPC** (Remote Procedure Call)

RPC: Basic idea

- Call a procedure, but execute it at a remote computer `res = proc(args)`
- Get the results *locally*, as if the call was *local*



- Request / reply mechanism
- Procedure call, but in a *disjoint address space*, on a *different host*



RPC: basic idea

- Needed to formalize a separate, but general, *procedure call protocol*
- Idea proposed by J. E. White at SRI in 1976
- Fully described by Birrell and Nelson in the early 1980's
 - Based on the work done at Xerox PARC

Andrew D. Birrell and Bruce Jay Nelson. 1984. *Implementing remote procedure calls*. ACM Trans. Comput. Syst. 2, 1 (February 1984), 39-59.



RPC: issues

- Issues to solve:
 - Binding (*connecting* the client and server)
 - Communication protocol
 - Dealing with failures – network/server crashes
 - Addressable arguments
 - Integration with existing systems
 - Data Integrity and security



RPC: principles

- A server defines its *protocol (interface)* using a **Protocol Specification Language**¹
- IDL² specifies the names, parameters, and types for all client-callable server procedures
- It uses a specific syntax to define the above elements
- A **stub compiler (rpcgen)** reads the protocol spec file and produces two stub procedures for each server procedure (for both the client and server side).

1. <http://tools.ietf.org/html/rfc1831.html>

2. Now, it is referred to as **Interface Definition Language (IDL)**



RPC: principles

- The server programmer implements the server procedures and links them with the server-side stubs
- The client programmer implements the client program and links it with the client-side stubs
- The stubs are responsible for managing all details of the remote communication between client and server



RPC: Protocol Example

```
struct param_in {  
    long factor;  
};  
struct result_out {  
    long coverage;  
};  
program SMART_PROG {  
    version SMART_VERS {  
        result_out SMARTPROC(param_in) = 1; /* procedure # */  
    } = 1; /* version # */  
} = 0x31230000; /* service id */
```



RPC: Protocol Example

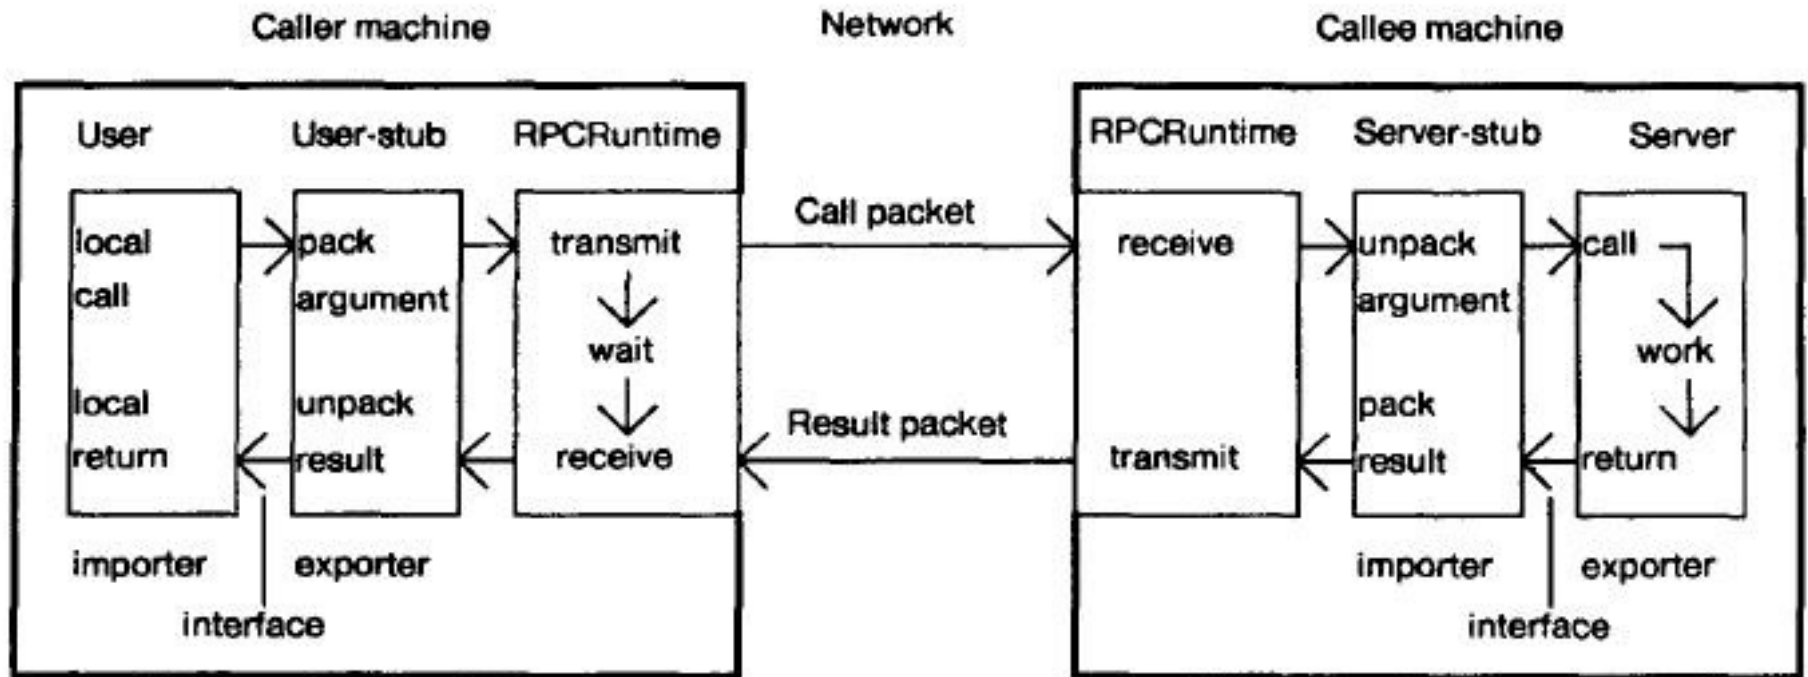
```
struct param_in {  
    long factor;  
};  
struct result_out {  
    long coverage;  
};  
program SMART_PROG {  
    version SMART_VERS {  
        result_out SMARTPROC(param_in) = 1; /* procedure # */  
    } = 1; /* version # */  
} = 0x31230000; /* service id */
```



RPC: operation

- A *client-side stub* is a procedure that looks to the client as if it were a callable (server) procedure
- A *server-side stub* looks to the server as if a (regular) client called it
- The client program thinks it is calling a “normal” C function; in fact, it is calling the client stub
- The server program thinks it is called from a C function; in fact, it’s called by the server stub
- The stubs send messages to each other to actually execute the call
- **RPC happens “transparently”**

RPC: operation





RPC: operation

- **Marshalling** is the packing of procedure (function) parameters into a message packet
- Basically, *what to do* (procedure) *with what* (parameters)
- The packet must be representable as a **sequence of bytes** to be written/read over a network
- The RPC stubs call type-specific procedures to *marshal* (or *unmarshal*) the parameters to a call



RPC: operation

- **Binding** is the process of connecting the client to the server
- The *server*, when it starts up, exports its interface
- Identifies itself to a network *name server*
- Tells RPC runtime it's alive and ready to accept calls
- Client uses the name server to find the location of a server and establish a connection
- Client is ready to invoke operations on the server



RPC: operation

- The client stub *marshals the operation and parameters* into a *message*
- The server stub *unmarshals operation and parameters* from the message and uses them to call the server operation (procedure) locally
- On return
 - The server stub *marshals the return value*
 - The client stub *unmarshals the return value* and returns it to the client program, as if it was produced locally



RPC: problems

- Issues with RPC
 - complexity
 - lack of language independent mechanisms (developed primarily for the C programming language)
- Performance overhead
 - often hand-tuned data transfer is more efficient
 - due to data encoding and to the overhead of authentication and fault tolerance



Distributed Objects

- Idea: extend the concepts of RPC to object-oriented style of programming
 - RPC for the “object crowd”
 - Remote method call
- Why Support Distributed Objects?
 - Object oriented paradigm
 - Hide implementation details under object interface
- Client code and object/class code potentially implemented in different OO languages



Distributed Objects

- CORBA (Common Object Request Broker Architect.)
 - Cross-lingual (primarily C++, Java)
 - Cross-platform
 - Many features
- DCOM (Distributed Component Object Model)
 - Microsoft's solution
 - Some cross-lingual support (within "Microsoft world")
 - Windows only
 - Built on DCE RPC and COM
- Java RMI
 - Single language, tightly integrated with Java



Java RMI

- Specification by Sun (now owned by Oracle) for making remote method invocations on Java objects from Java clients
- Specific to Java
- Interfaces and Classes for specification are
- Part of the core Java platform since 1.1 (1997)
- Java Standard Edition (J2SE) comes with the implementation of RMI



Java RMI: Marshalling

Serialization

- Conversion of a Java Object into a byte stream
- Any Java class instance can be *serialized*
- Involves writing and reading state of an object in the same order and using the same encoding
- Serialization can be
 - Provided automatically by the Java environment
 - Written explicitly by the developer writing the class
- Serialization is the “wire” protocol for Java RMI

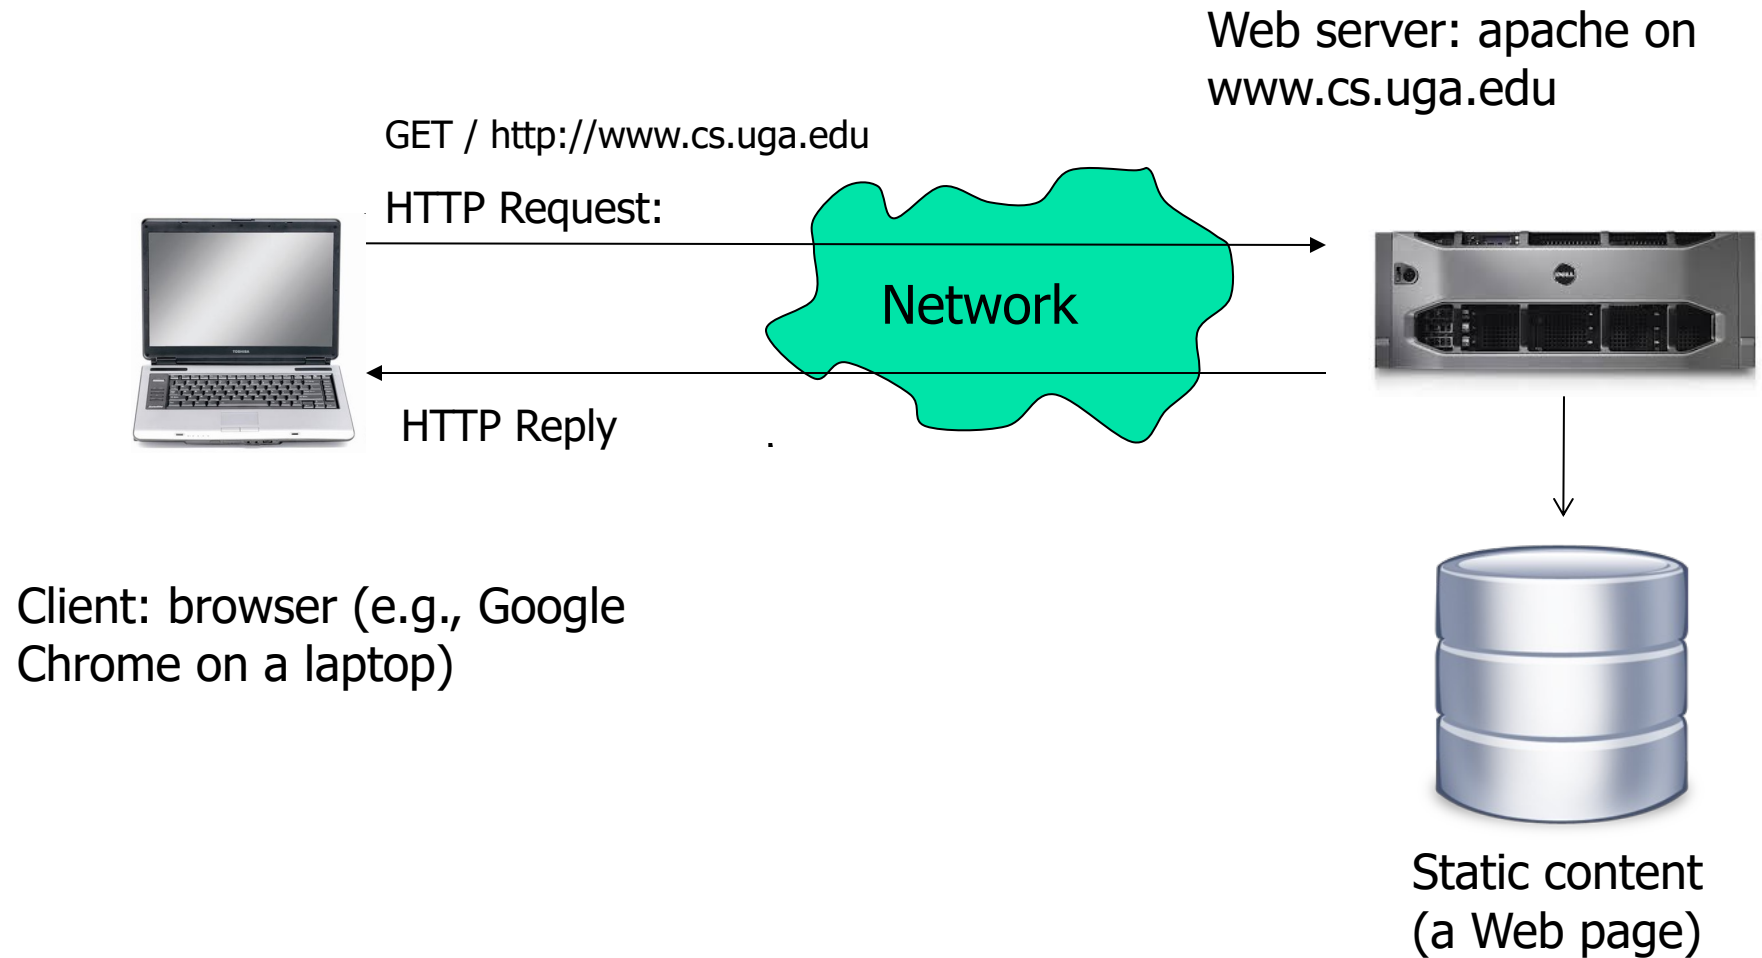


World Wide Web

World Wide Web

- Revolutionized remote access to data
- Data, referred to as ***resources***, can be remotely requested, downloaded/uploaded
- Resources are delivered by Web servers, or similar systems
- HTTP (Hypertext Transfer Protocol) is used to communicate between the client (requestor) and the server

WWW communication



Web communication

Client


connect to `www.cs.uga.edu`

Server

accept connection

GET / HTTP/1.0

Language agnostic



file content

Retrieve a file
(resource)
from disk



HTTP Request

- A request (HTTP 1.0) has the following form:

```
Request-Method Request-URL HTTP-Version <CR><LF>
(generic-header | request-header | entity-header) <CR><LF>
<CR><LF>
[message body]
```

- For example, on cobweb.cs.uga.edu (using HTTP 1.0):

GET / HTTP/1.0

- A simple request to www.cs.uga.edu (using HTTP 1.1):

GET / HTTP/1.1

Host: www.cs.uga.edu



HTTP Request

- HTTP request methods:

- **GET** – request whatever information is identified by the Request-URL
- **POST** – request that server accepts the entity enclosed in the request (create a child resource)
- **PUT** – request that the enclosed entity be stored under the Request-URL (replace an existing or create new)
- **DELETE** – request that the server delete the resource identified by Request-URL
- **HEAD** – identical to GET, but server must not return a message body in response
- A few additional methods are available




Dynamic Content

- Responses (usually, Web pages) can be created **dynamically**, in response to requests
- Advantages
 - personalization (e.g., athena.uga.edu),
 - interaction with client input
 - interaction with back-end applications
- Disadvantages
 - performance penalty -- dynamic content must be **created at request time**, which takes time
- Generating dynamic content (CGI Scripts, Servlets)

Web communication

Web server: Apache on
uml.cs.uga.edu



Client: browser (e.g., Google
Chrome on a laptop)

CGI

```
17 string sInput;  
18 int iLength, iN;  
19 double dblTemp;  
20 bool again = true;  
21  
22 while (again) {  
23     iN = -1;  
24     again = false;  
25     getline(cin, sInput);  
26     stringstream(sInput) >> dblTemp;  
27     iLength = sInput.length();  
28     if (iLength < 4) {  
29         again = true;  
30         continue;  
31     } else if (sInput[iLength - 3] != '.') {  
32         again = true;  
33         continue;  
34     } while (++iN < iLength) {  
35         if (isdigit(sInput[iN])) {  
36             continue;  
37         } else if (iN == (iLength - 3)) {  
38             continue;  
39         }  
40     }  
41     // ...  
42 }
```

CGI is Common Gateway Interface

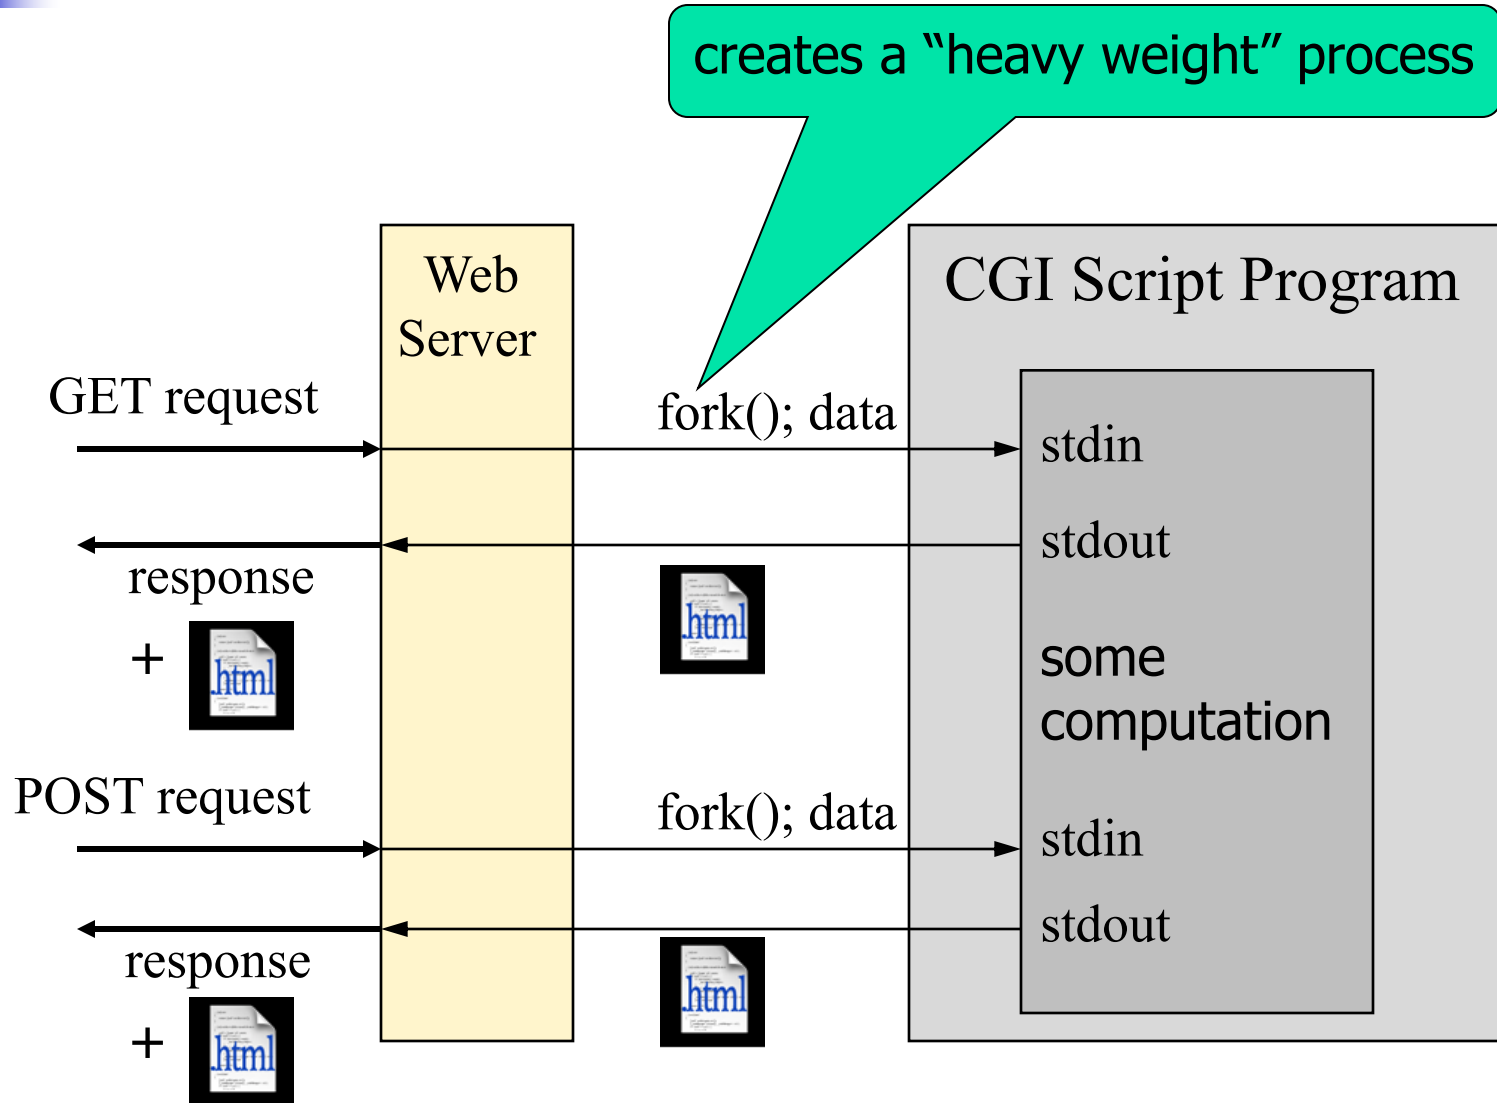
Dynamic content
(a CGI script)



CGI Scripts

- CGI scripts were introduced to provide dynamic content by **executing programs remotely**
- The script is a program (e.g., C, PHP, Perl)
- A client requests a dynamic resource (URL)
- When the URL is requested, the **server invokes the named script**, passing to it client data (input parameters)
- The **script** executes and ultimately **outputs an HTML page** which is sent back to the client

CGI Execution





Servlets

- When we run small Java programs within a browser these are referred to as *applets*
- Small Java programs running within a server are called *servlets*
- A *servlet* is a server-side program designed to process a client request (which requires interactivity)
- Servlets have become a popular mechanism for creating *interactive* servers

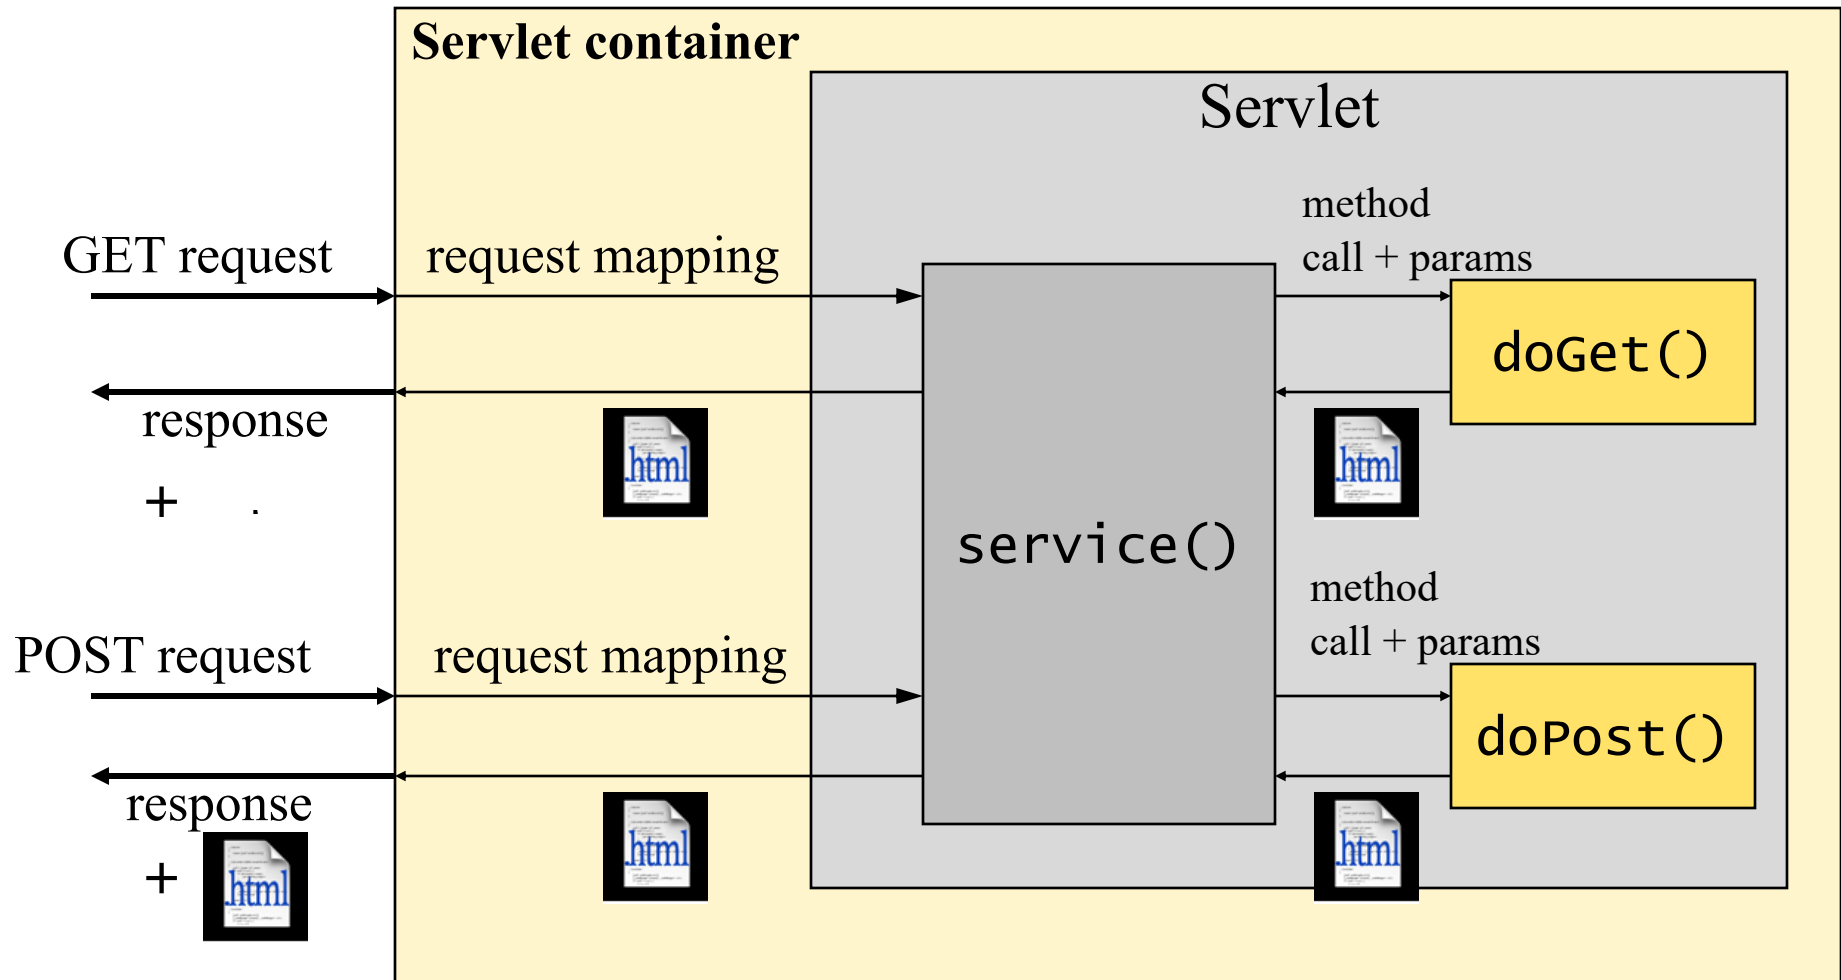


Servlets

- A servlet (an instance of a special Java class) is compiled and deployed to run within a special server
- Such a server is frequently called a **servlet container**
- A servlet container with greatly enhanced capabilities (messaging, administration, etc.) is referred to as an **application server**

Servlet invocation

Once started, a servlet does not exit





Web Services

- Consider a servlet container with many servlets, each for some type of a remote computation to produce desired data
- Collectively, these servlets can be thought of as a **service** providing some programs (services) to many clients
- **Web services** have been introduced as a generalization of server-side WWW processing



Web Services

- Web Services – the name
 - Hewlett-Packard's e-Speak in 1999, an enabler for e-services
 - Microsoft introduced the name "Web services" in June 2000
 - Now, Web services have become virtually ubiquitous



Web Services

- Two basic types of Web Services:
 - SOAP Web Services
 - REST Web Services
- The same (old) problem: how to invoke a remote computation over the network and get the result?
- We will very briefly talk about SOAP and then focus on REST services



Web Services

SOAP (Simple Object Access Protocol) is Based on open, standard technologies

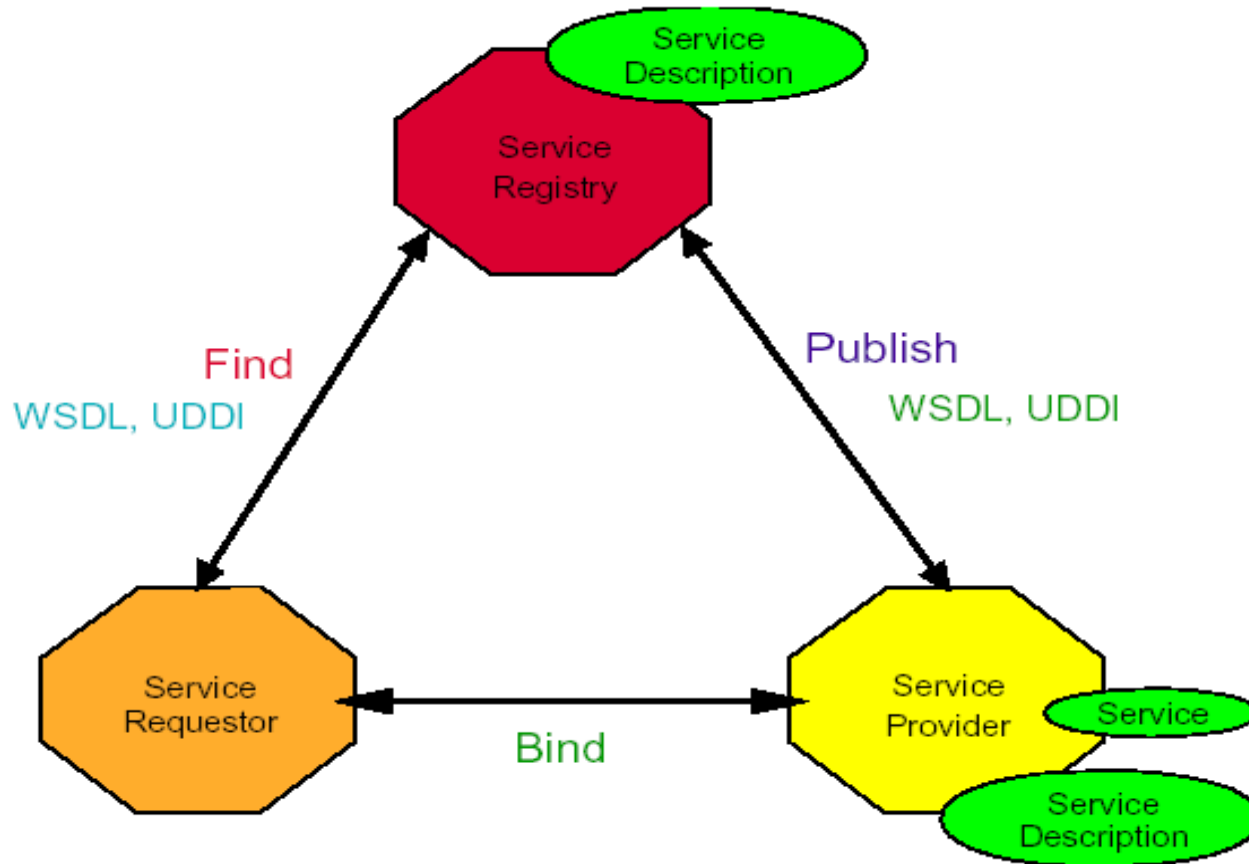
- XML – encoding data such that it can be exchanged between applications and platforms
- SOAP is a messaging protocol for transporting information and instructions between applications (uses XML)



Web Services

- WSDL (Web Services Description Language)
a standard method of describing Web Services and their capabilities (in XML)
- UDDI (Universal Description, Discovery and Integration)
defines XML-based rules for building directories in which service providers advertise their Web Services.

Web Services: Model

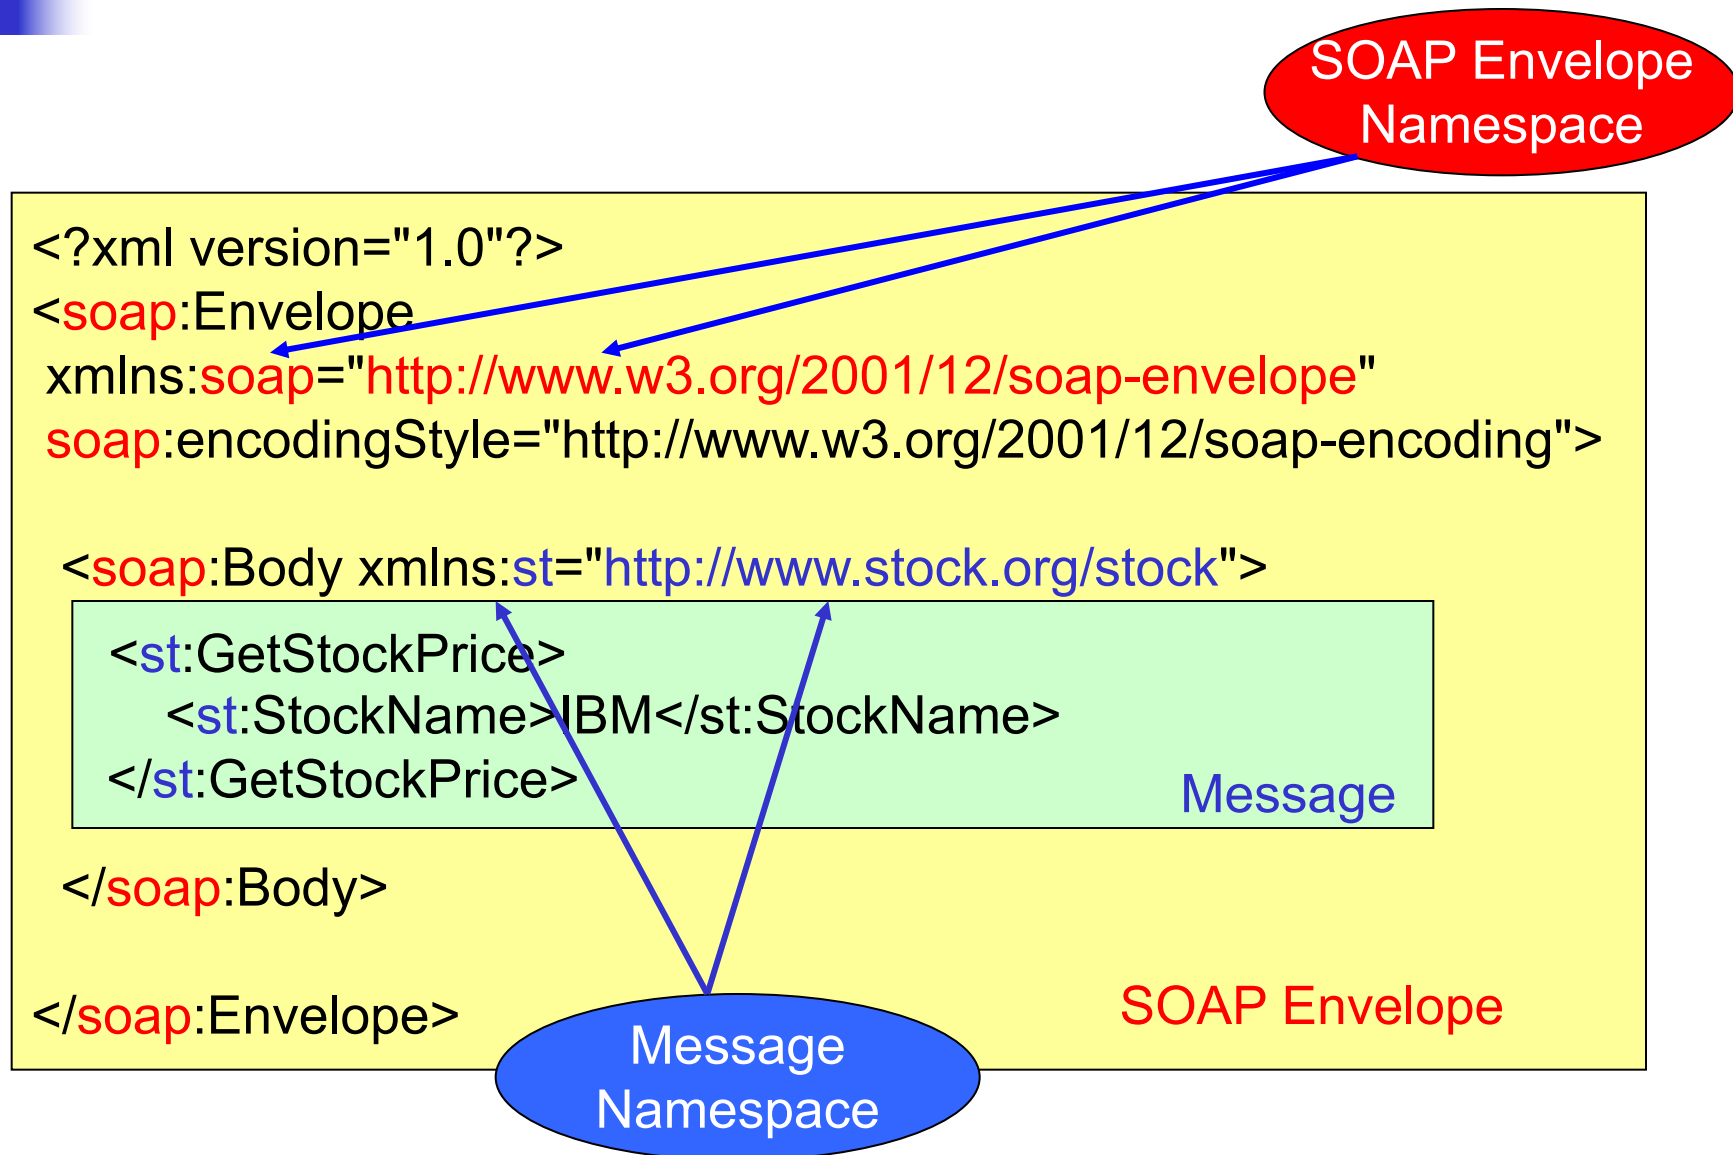




Web Services: SOAP

- Platform and language independent
- SOAP message has three parts
 - *envelope* – wraps entire message and contains header and body
 - *header* – optional element with additional info such as security or routing
 - *body* – application-specific data being communicated

SOAP Request Message





SOAP Response Message

```
<?xml version="1.0"?>
<soap:Envelope
xmlns:soap="http://www.w3.org/2001/12/soap-envelope"
soap:encodingStyle="http://www.w3.org/2001/12/soap-encoding">
  <soap:Body xmlns:st="http://www.stock.org/stock">
```

```
    <st:GetStockPriceResponse>
      <st:Price>34.5</st:Price>
    </st:GetStockPriceResponse>
```

Message

```
</soap:Body>
</soap:Envelope>
```

SOAP Envelope

Result
returned in
Body





SOAP Services: Verb-centered

- Consider the remote calls considered so far:
 - RPC (Remote Procedure Call)
 - RMI (Remote Method Invocation)
 - Even a plain Java or C++ method call
- They all share a common aspect: their names correspond to **actions**, usually denoted as **verbs** or **verb phrases**:

setText, compute, solve, draw, print



REST Services: Noun-centered

- We will transition to Web services where the focus is on **resources** (**nouns**), instead of actions (verbs)
- The actions are *fixed*
- Actions correspond to the main HTTP methods: GET, POST, PUT, and DELETE